

Distributed Systems
Master of Science in Engineering in Computer
Science

AA 2020/2021

LECTURE 3 (PART 2): TIME IN DISTRIBUTED SYSTEMS

Logical Time

Logical clock

Physical clock synchronization algorithms try to coordinate distributed clocks to reach a common value

- Physical clock synchronization algorithms are based on the estimation of transmission delay but in several system it can be hard to find a good estimation.
- In several applications it is not important when things happened but in which order they happened

However in a Distributed System, each system has its own “logical clock”

- If clocks are not aligned it is not possible to order events generated by different processes

Reliable way of ordering events is required!

■ **Notes:**

- Two events occurred at some process p_i happened in the same order as p_i observes them
 - When p_i sends a message to p_j the *send* event happens before the *receive* event
- Lamport introduces the *happened-before* relation that captures the causal dependencies between events (*causal order relation*)
- We note with \rightarrow_i the ordering relation between events in a process p_i
 - We note with \rightarrow the happened-before between any pair of events

Happened-Before Relation: Definition

Two events e and e' are related by happened-before relation ($e \rightarrow e'$) if:

- $\exists p_i \mid e \rightarrow_i e'$
- $\forall \text{ message } m \text{ send}(m) \rightarrow \text{receive}(m)$
 - $\text{send}(m)$ is the event of sending a message m
 - $\text{receive}(m)$ is the event of receipt of the same message m
- $\exists e, e', e'' \mid (e \rightarrow e'') \wedge (e'' \rightarrow e')$ (happened-before relation is transitive)


Happened-Before Relation

Using the three rules is possible to define a causal ordered sequence of events e_1, e_2, \dots, e_n

Notes:

- The sequence e_1, e_2, \dots, e_n may not be unique
- It may exist a pair of events $\langle e_1, e_2 \rangle$ such that e_1 and e_2 are not in happened-before relation
- If e_1 and e_2 are not in happened-before relation then they are *concurrent* ($e_1 \parallel e_2$)
- For any two events e_1 and e_2 in a distributed system, either $e_1 \rightarrow e_2$, $e_2 \rightarrow e_1$ or $e_1 \parallel e_2$

happened-before: example



e^j_i is j -th event of process p_i

$$S_1 = \langle e^1_1, e^1_2, e^2_2, e^2_3, e^3_3, e^3_1, e^4_1, e^5_1, e^4_2 \rangle$$

$$S_2 = \langle e^1_3, e^2_1, e^3_1, e^4_1, e^5_3 \rangle$$

Note:

e^1_3 and e^1_2 are concurrent

Logical Clock

The Logical Clock, introduced by Lamport, is a software counting register *monotonically* increasing its value

- Logical clock is not related to physical clock

Each process p_i employs its logical clock L_i to apply a *timestamp* to events

$L_i(e)$ is the “logical” timestamp assigned, using the logical clock, by a process p_i to events e .

Property:

- If $e \rightarrow e'$ then $L(e) < L(e')$

Observation:

- The ordering relation obtained through logical timestamps is only a partial order. Consequently timestamps could not be sufficient to relate two events

Scalar Logical Clock: an implementation

Each process p_i initializes its logical clock $L_i=0$ ($\forall i = 1....N$)

p_i increases L_i of 1 when it generates an event (either *send* or *receive*)

- $L_i = L_i + 1$


When p_i sends a message m

- creates an event *send*(m)
- increases L_i
- timestamps m with $t=L_i$

When p_i receives a message m with timestamp t

- Updates its logical clock $L_i = \max(t, L_i)$
- Produces an event *receive*(m)
- Increases L_i

Scalar Logical Clock: example



- e_i^j is j -th event of process p_i
- L_i is the logical clock of p_i
- Note:
 - $e_1^1 \rightarrow e_2^1$ and timestamps reflect this property
 - $e_1^1 \parallel e_3^1$ and respective timestamps have the same value
 - $e_2^1 \parallel e_3^1$ but respective timestamps have different values

Limits of Scalar Logical Clock

Scalar logical clock can guarantee the following property

- IF $e \rightarrow e'$ then $L(e) < L(e')$

But it is not possible to guarantee

- IF $L(e) < L(e')$ then $e \rightarrow e'$

Consequently:

- It is not possible to determine, analyzing only scalar clocks, if two events are concurrent or correlated by the happened-before relation.

Mattern [1989] and Fridge [1991] proposed an improved version of logical clock where events are timestamped with local logical clock and node identifier

- ***Vector Clock***

Vector Clock : definition

Vector Clock for a set of N processes is composed by an array of N integer counters

Each process p_i maintains a Vector Clock V_i and timestamps events by mean of its Vector Clock

Similarly to scalar clock, Vector Clock is attached to message m (in this case we attach an array of integer)

Vector Clock allows nodes to order events in happens-before order based on timestamps

- Scalar clocks: $e \rightarrow e'$ implies $L(e) < L(e')$
- Vector clocks: $e \rightarrow e'$ **iff** $L(e) < L(e')$

Vector Clock : an implementation

Each process p_i initializes its Vector Clock V_i

- $V_i[j] = 0 \quad \forall j = 1 \dots N$

p_i increases $V_i[i]$ of 1 when it generates an event

- $V_i[i] = V_i[i] + 1$

When p_i sends a message m

- Creates an event *send*(m)
- Increases V_i
- timestamps m with $t = V_i$

When p_i receives a message m containing timestamp t

- Updates its logical clock $V_i[j] = \max(t[j], V_i[j]) \quad \forall j = 1 \dots N$
- Generates an event *receive*(m)
- Increases V_i

Vector Clock: properties

A Vector Clock V_i

- $V_i[i]$ represents the number of events produced by p_i
- $V_i[j]$ with $i \neq j$ represents the number of events generated by p_j that p_i can know

$V = V'$ if and only if

- $V[j] = V'[j] \quad \forall j = 1 \dots N$


$V \leq V'$ if and only if

- $V[j] \leq V'[j] \quad \forall j = 1 \dots N$

$V < V'$ therefore the event associated to V happened before the event associated to V' if and only if

- $V \leq V' \wedge V \neq V'$
 - $\forall i = 1 \dots N \quad V'[i] \geq V[i]$
 - $\exists i \in \{1 \dots N\} \mid V'[i] > V[i]$

Vector Clock: an example



A comparison of Vector Clocks

1
0
0

V

1
1
0

V'

$V(e) < V'(e')$ then $e \rightarrow e'$

1
0
0

V

0
0
1

V'

$V(e) \neq V'(e')$ then $e \parallel e'$

Differently from Scalar Clock, Vector Clock allows to determine if two events are concurrent or related by an happened-before relation

Logical Time and Distributed Algorithms

Logical clock in distributed algorithms

We have seen two mechanisms to represent logical time

- Scalar Clock
- Vector Clock


Each mechanism can be used to solve different problems, depending on the problem specification

- Scalar Timestamp → Lamport's Mutual Exclusion
- Vector Timestamp → Causal Broadcast

The Mutual Exclusion abstraction

Specification

- **Mutual Exclusion**: at every time t at most one process p is in critical section
- **No-Deadlock**: there always exists a process p able to enter the critical section.
- **No-Starvation**: every process p requesting the critical section eventually gets in.



Time stamp based algorithm: Lamport

Difference from concurrent system

- When a process wants to enter the CS sends a request message to all the other

An history of the operations is maintained by using a counter (time stamp)

Each transmission and reception event is relevant to the computation:

- The counter is incremented for each send and receive event
- The counter is incremented also when a message, not directly related to the mutual exclusion computation, is sent or received.

Lamport's algorithm: implementation

Local data structures to each process p_i

- ck
 - Is the counter for process p_i
- Q
 - Is a queue maintained by p_i where CS access requests are stored

Algorithm rules for a process p_i


- Access the CS
 - p_i sends a request message, attaching ck , to all the other processes
 - p_i adds its request to Q
- Request reception from a process p_j
 - p_i puts p_j request (including the timestamp) in its queue
 - p_i sends back an ack to p_j

Lamport's algorithm: implementation

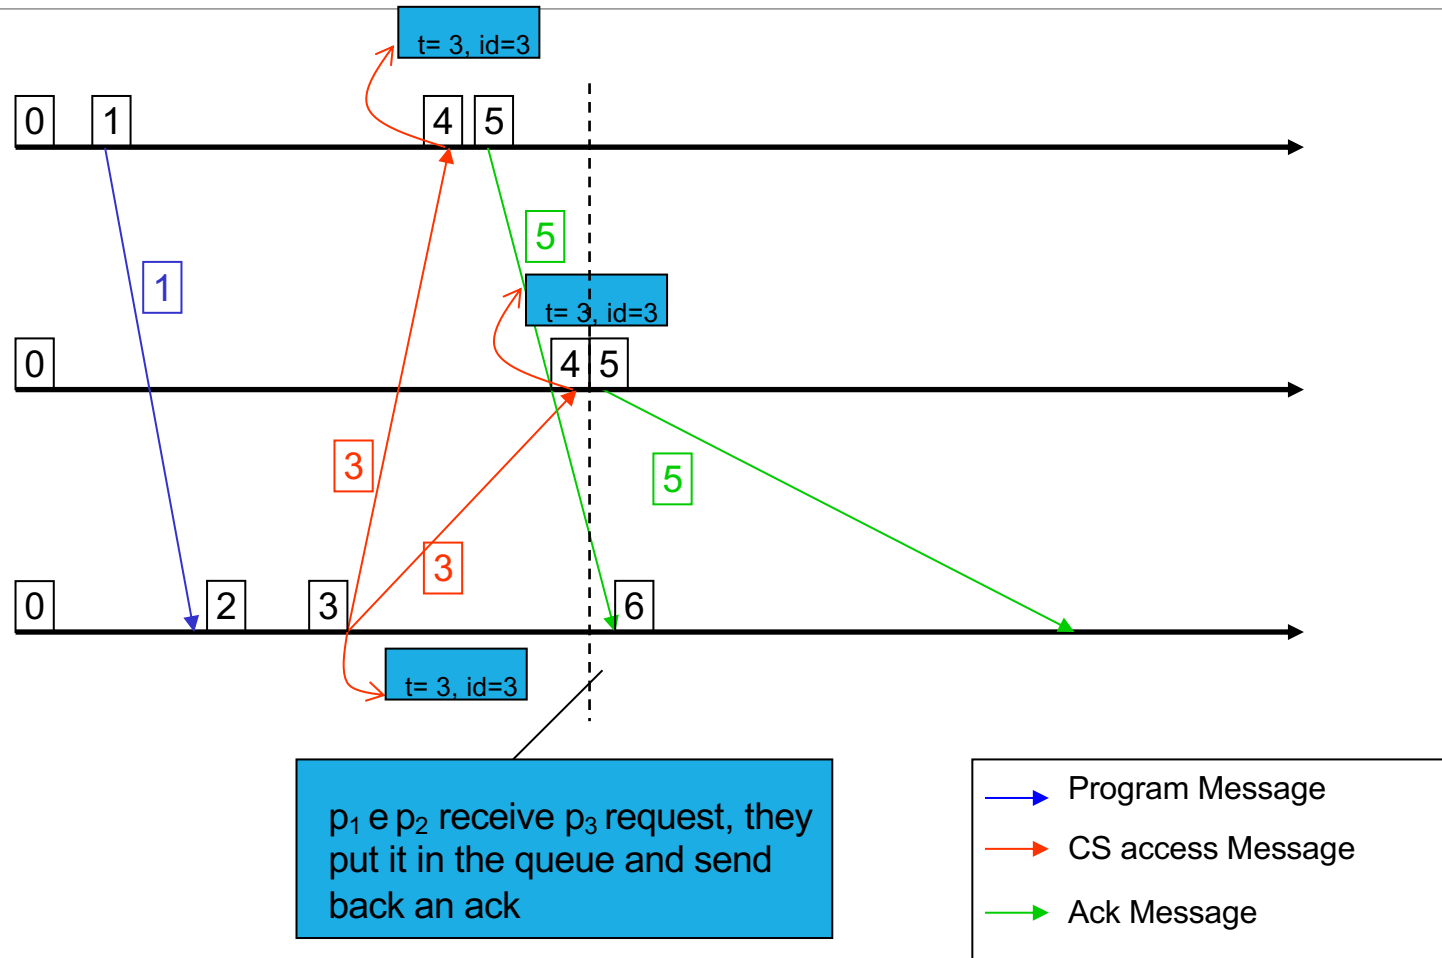
Algorithm rules for a process p_i

- p_i enters the CS iff
 - p_i has, in its queue, a request with timestamp t
 - t is the small timestamp in the queue
 - p_i has already received an ack with timestamp t' from any other process and $t' > t$
- Release of the CS
 - p_i sends a RELEASE message to all the other processes
 - p_i deletes its request from the queue
- Reception of a release message from a process p_j
 - p_i deletes p_j 's request from the queue

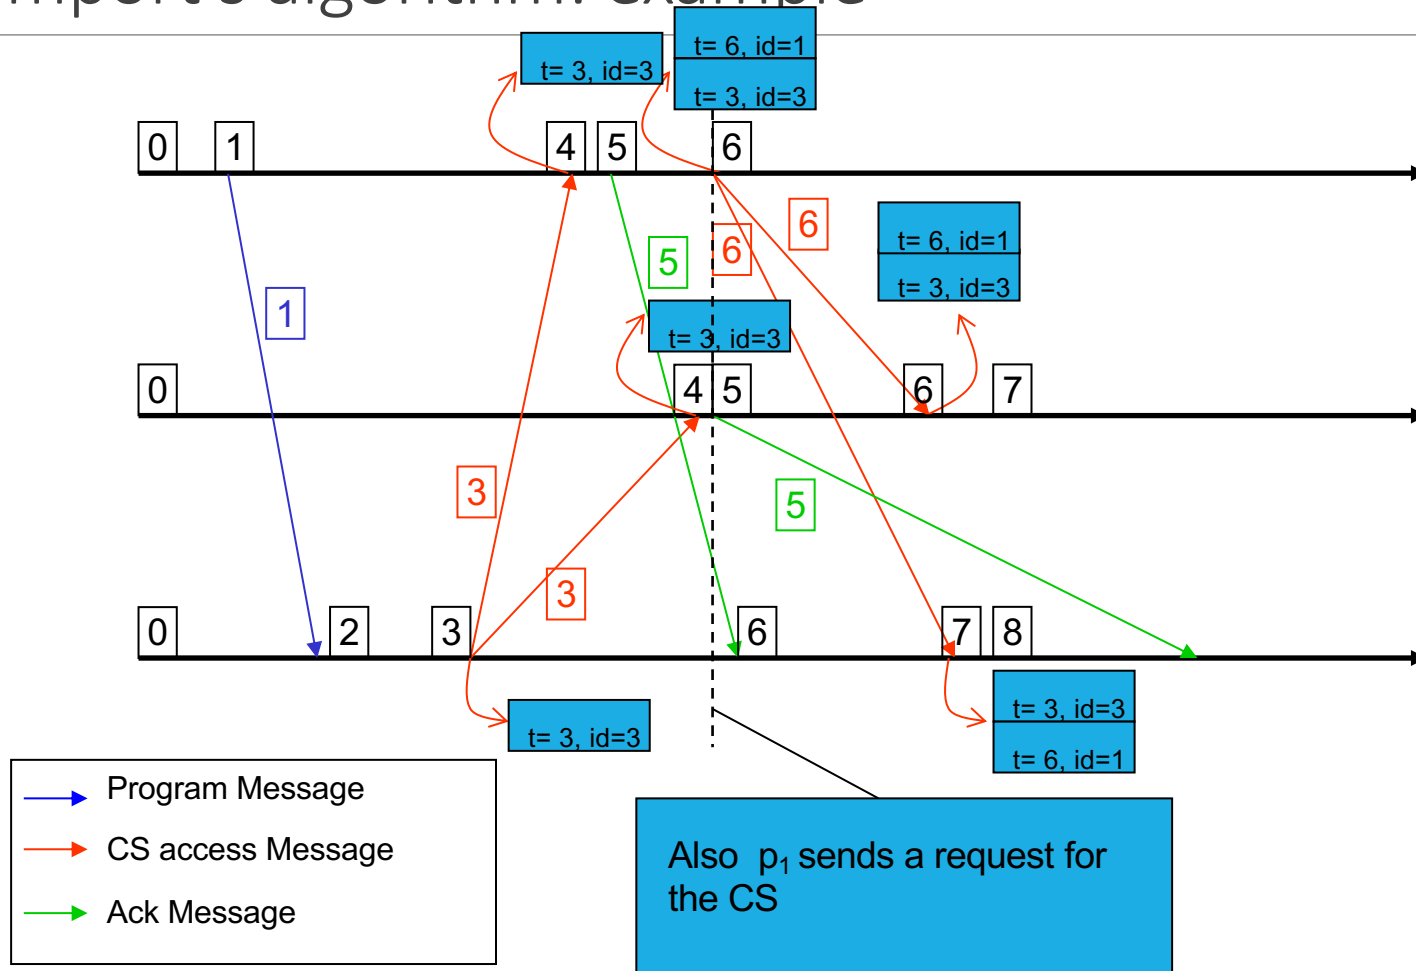
Lamport's algorithm: example



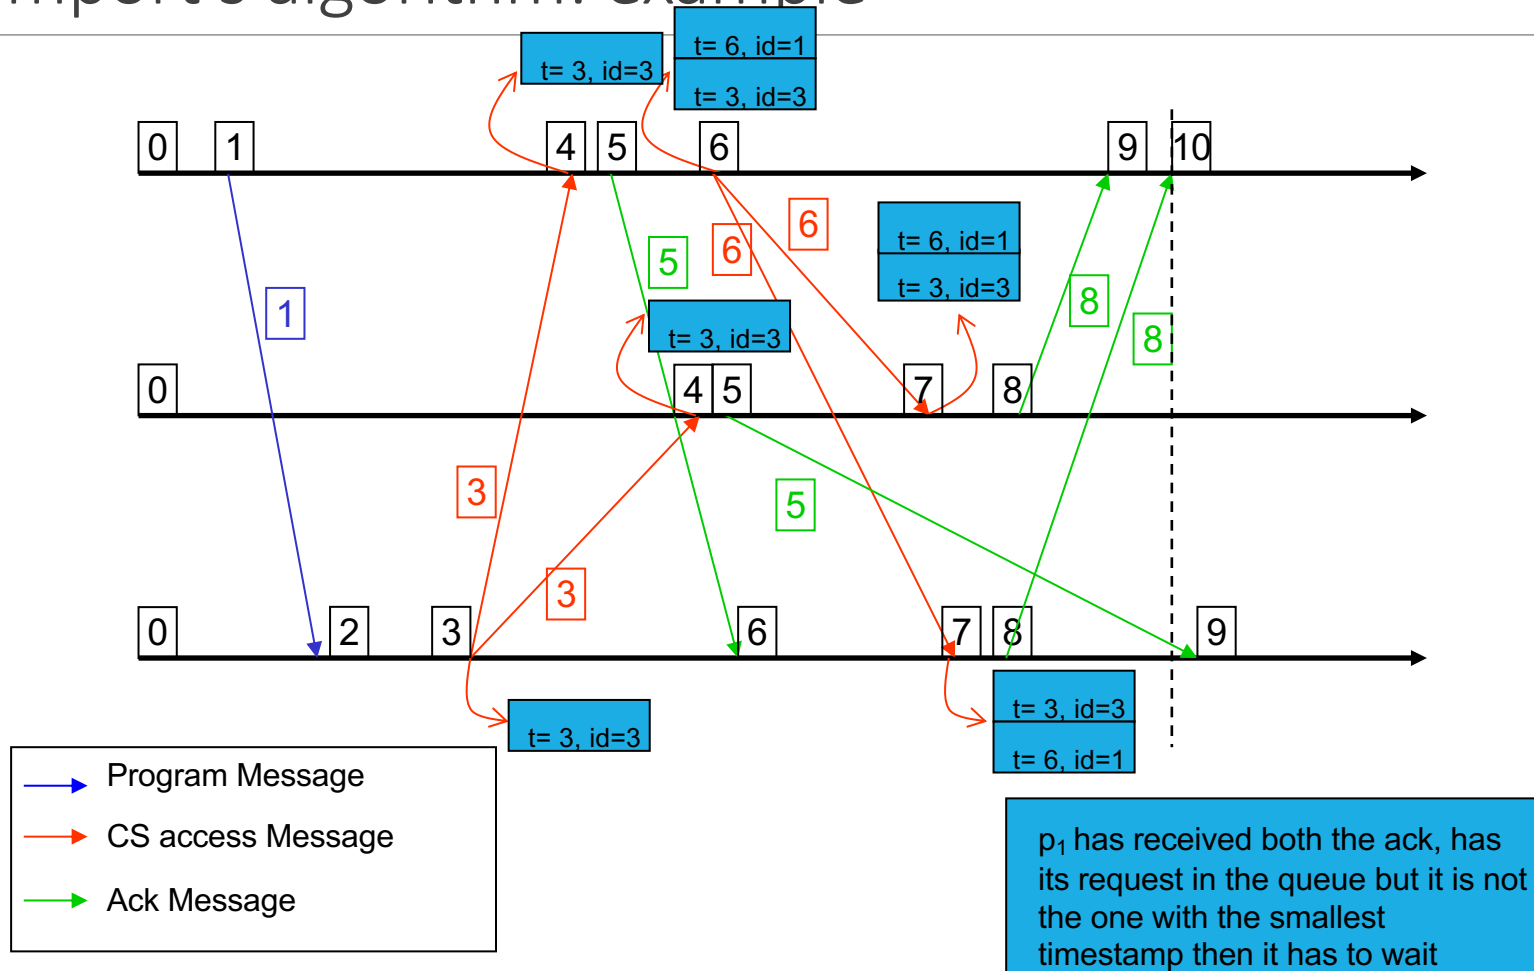
Lamport's algorithm: example



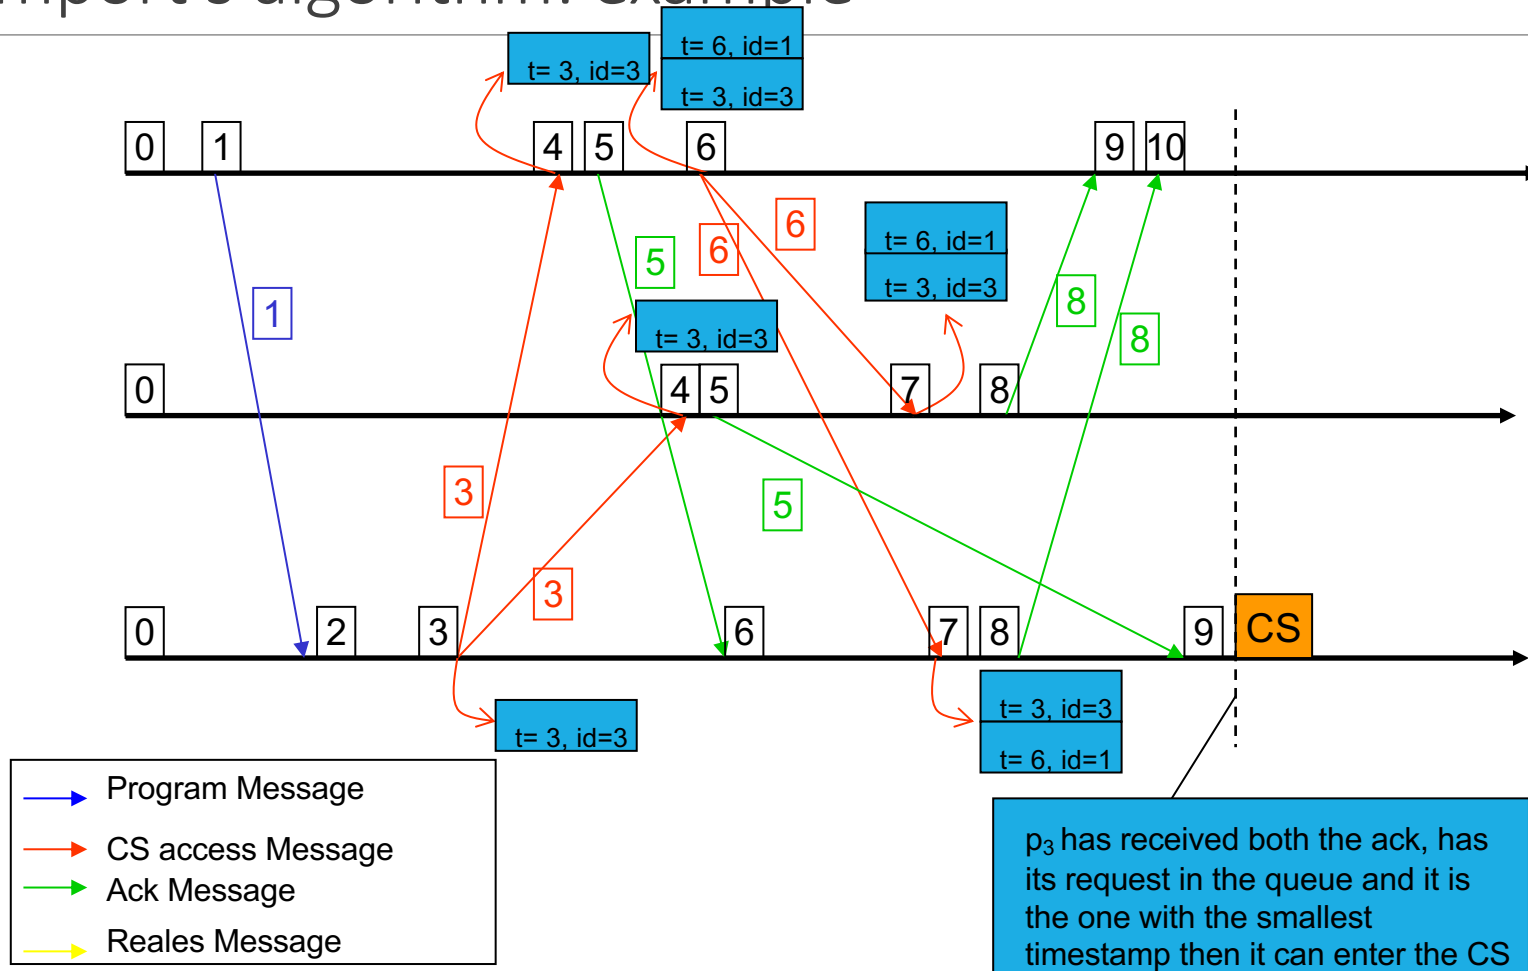
Lamport's algorithm: example



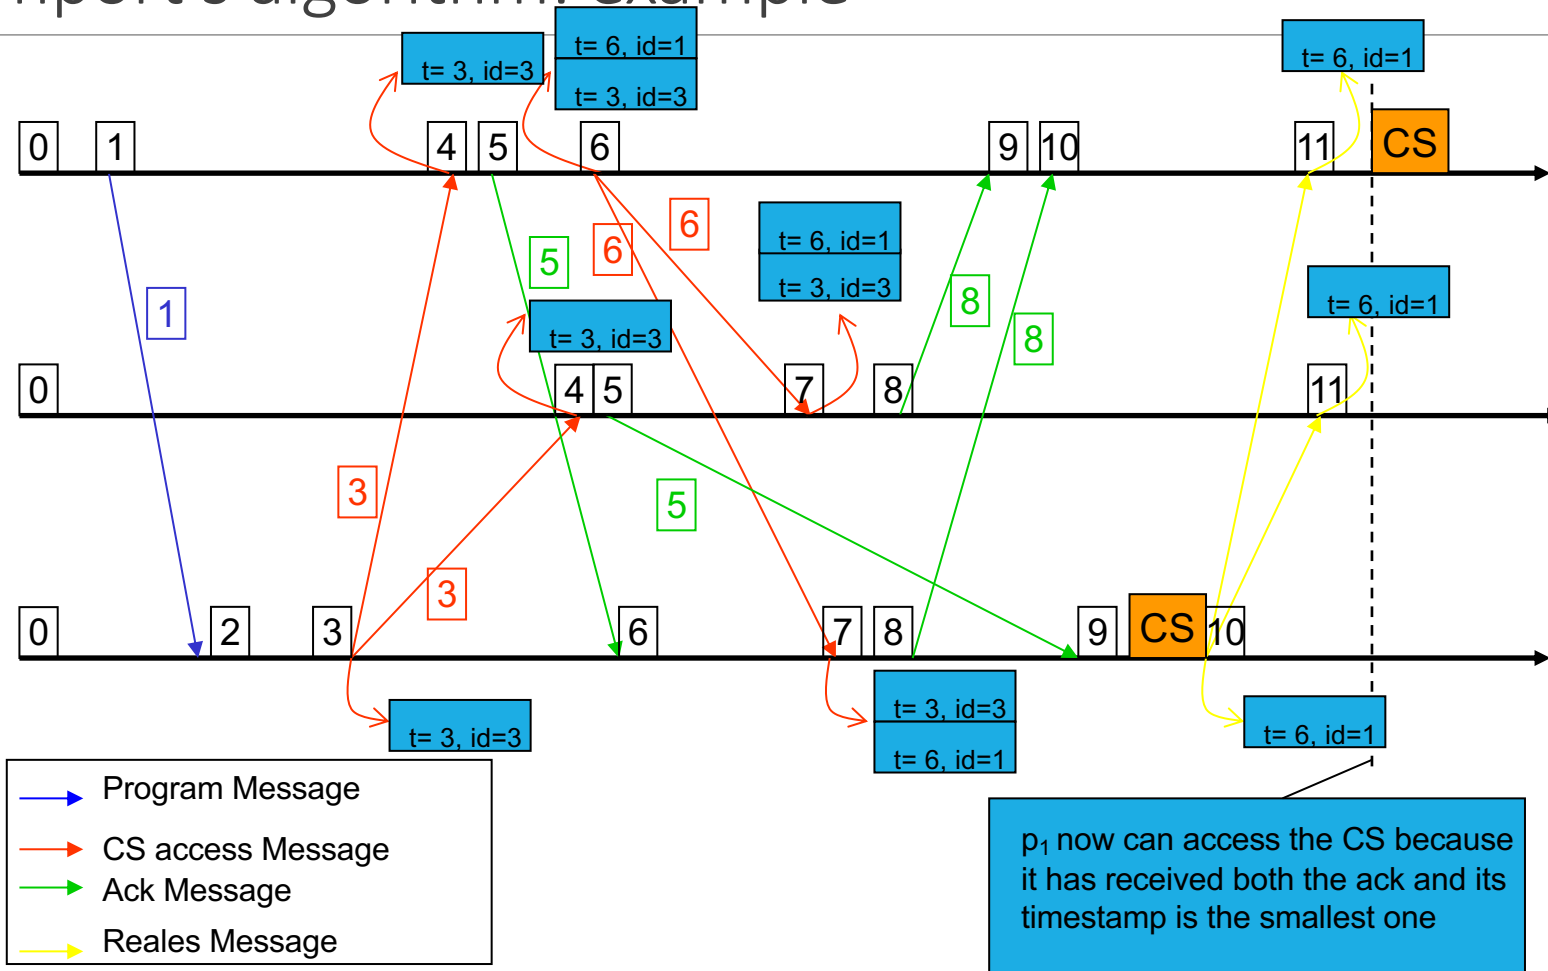
Lamport's algorithm: example



Lamport's algorithm: example




Lamport's algorithm: example



Lamport's algorithm: safety proof

Let us suppose by contradiction that both p_i and p_j enter the CS

- \Rightarrow both the processes have received an ack from any other process and, to enter the CS, the timestamp has to be the smallest in the queue
 - $t_i < t_j < \text{ack}_i.\text{ts}$
 - $t_j < t_i < \text{ack}_j.\text{ts}$



Lamport's algorithm: properties

Fairness is satisfied: different requests are satisfied in the same order as they are generated

- Such order comes from the happened-before relation:
 - If two requests are in happened-before relation then they are satisfied in the same order.
 - If two request are concurrent with respect to the happened before relation then the access can happen in any order

Lamport's algorithm: performances

Lamport's algorithm needs $3(N-1)$ messages for the CS execution

- $N-1$ requests
- $N-1$ acks
- $N-1$ releases

In the best case (none is in the CS and only one process ask for the CS) there is a delay (from the request to the access) of 2 messages

Ricart-Agrawala's algorithm: implementation

Local variables

- #replies (initially 0)
- State $\in \{\text{Requesting}, \text{CS}, \text{NCS}\}$ (initially NCS)
- Q pending requests queue (initially empty)
- Last_Req
- Num

Algorithm

begin

1. State=Requesting
2. Num=num+1; Last_Req=num
3. $\forall i=1\dots N$ send REQUEST(num) to p_i
4. Wait until #replies=n-1
5. State=CS
6. CS
7. $\forall r \in Q$ send REPLY to r
8. $Q = \emptyset$; State=NCS; #replies=0

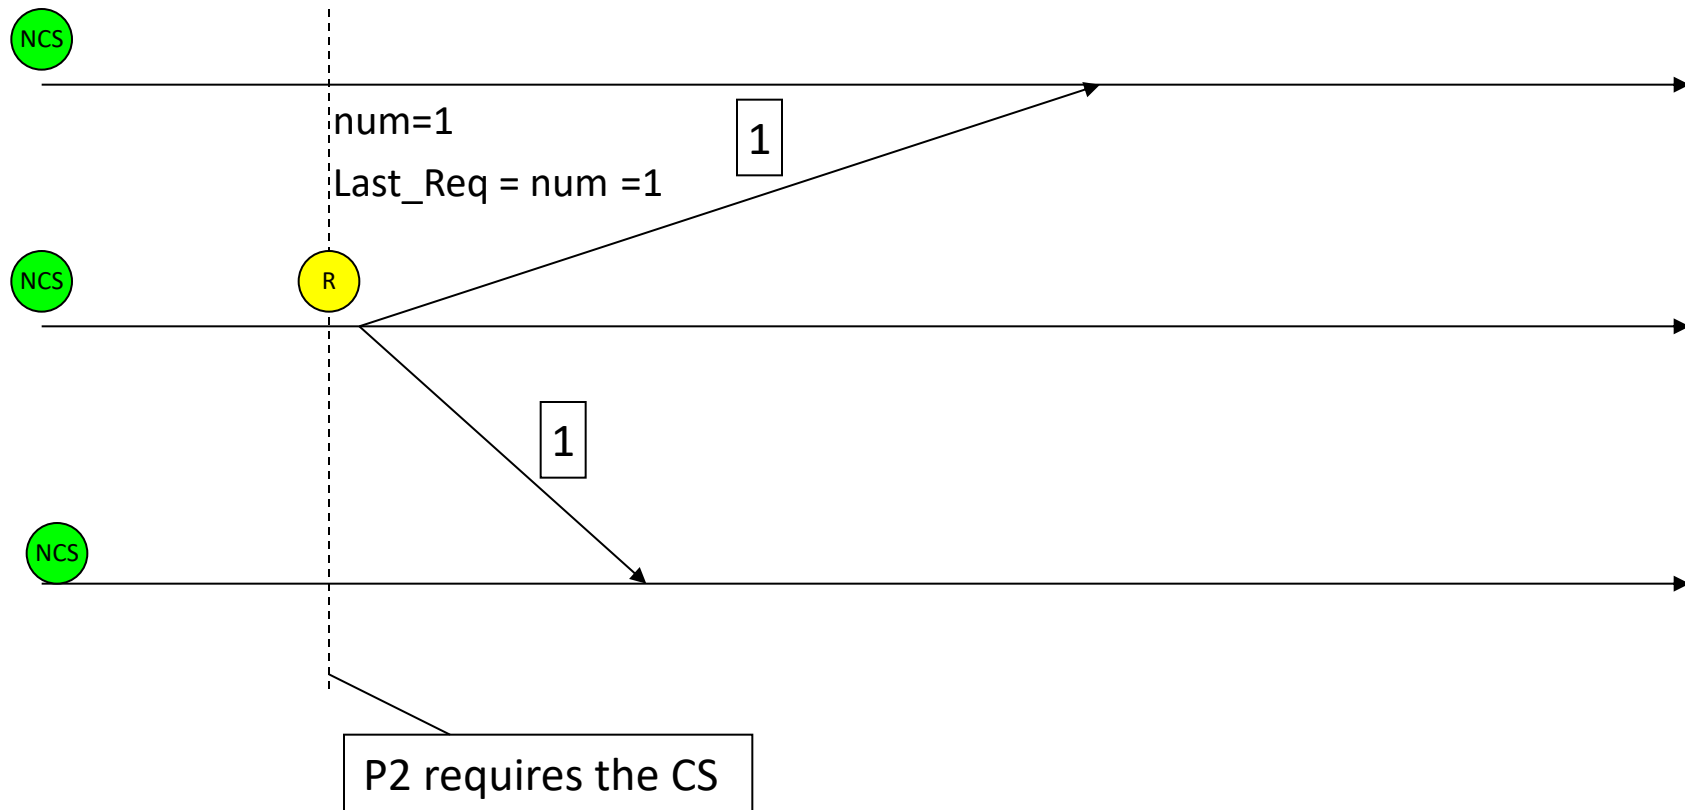
Upon receipt REQUEST(t) from p_j

1. If State=CS or (State=Requesting and $\{\text{Last_Req}, i\} < \{t, j\}$)
2. Then insert in $Q\{t, j\}$
3. Else send REPLY to p_j
4. Num=max(t,num)

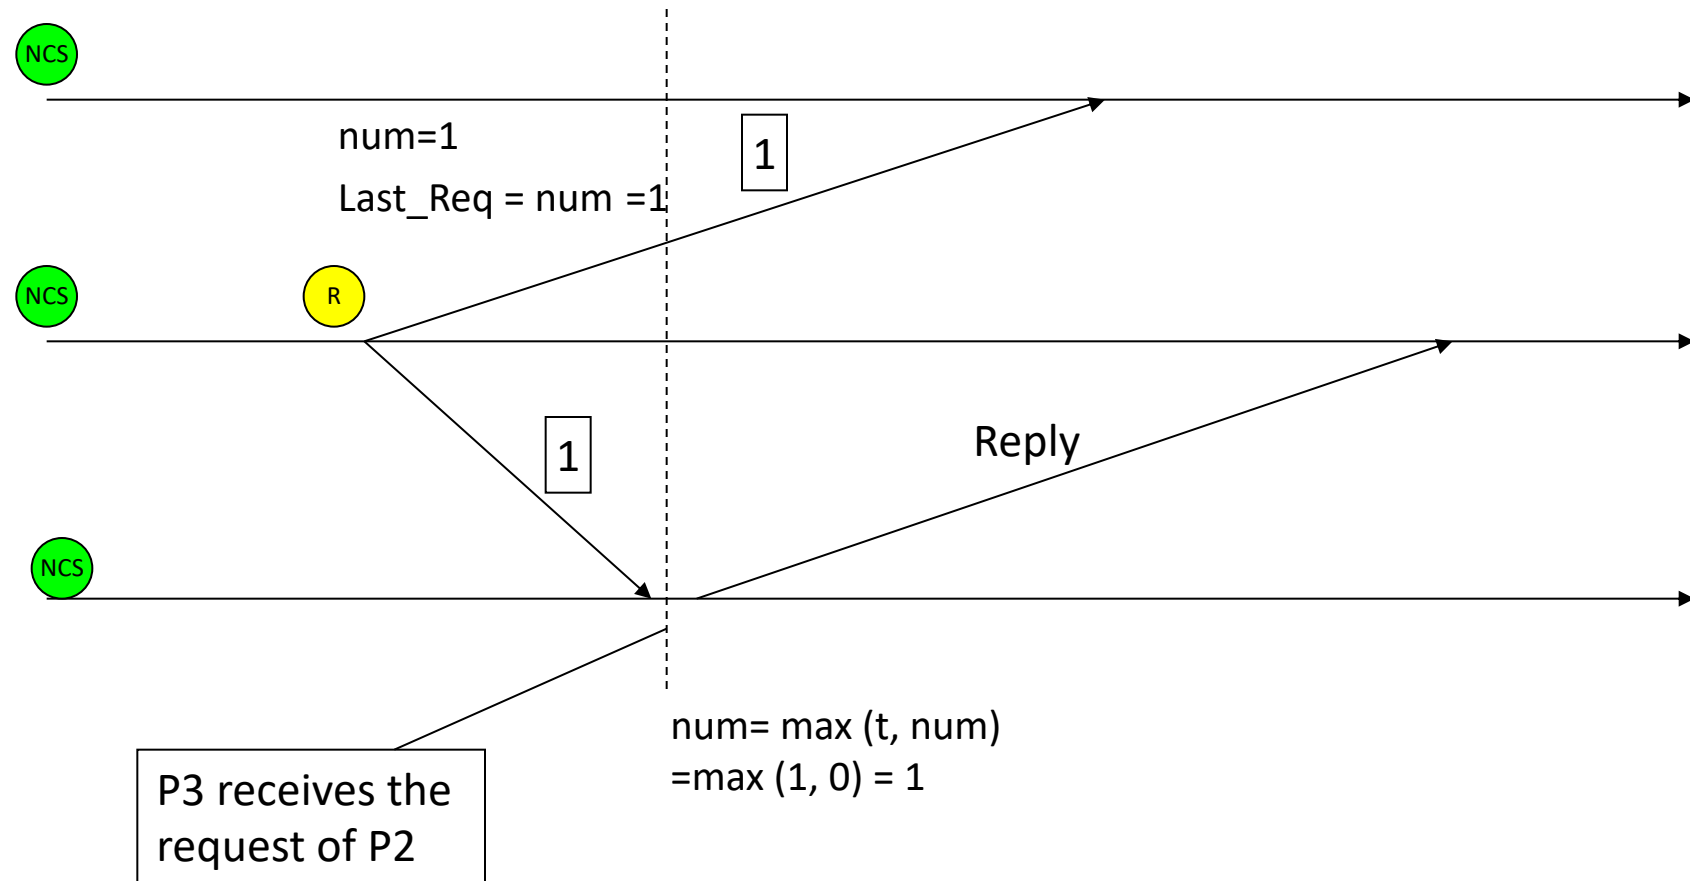
Upon receipt of REPLY from p_j

1. #replies=#replies+1

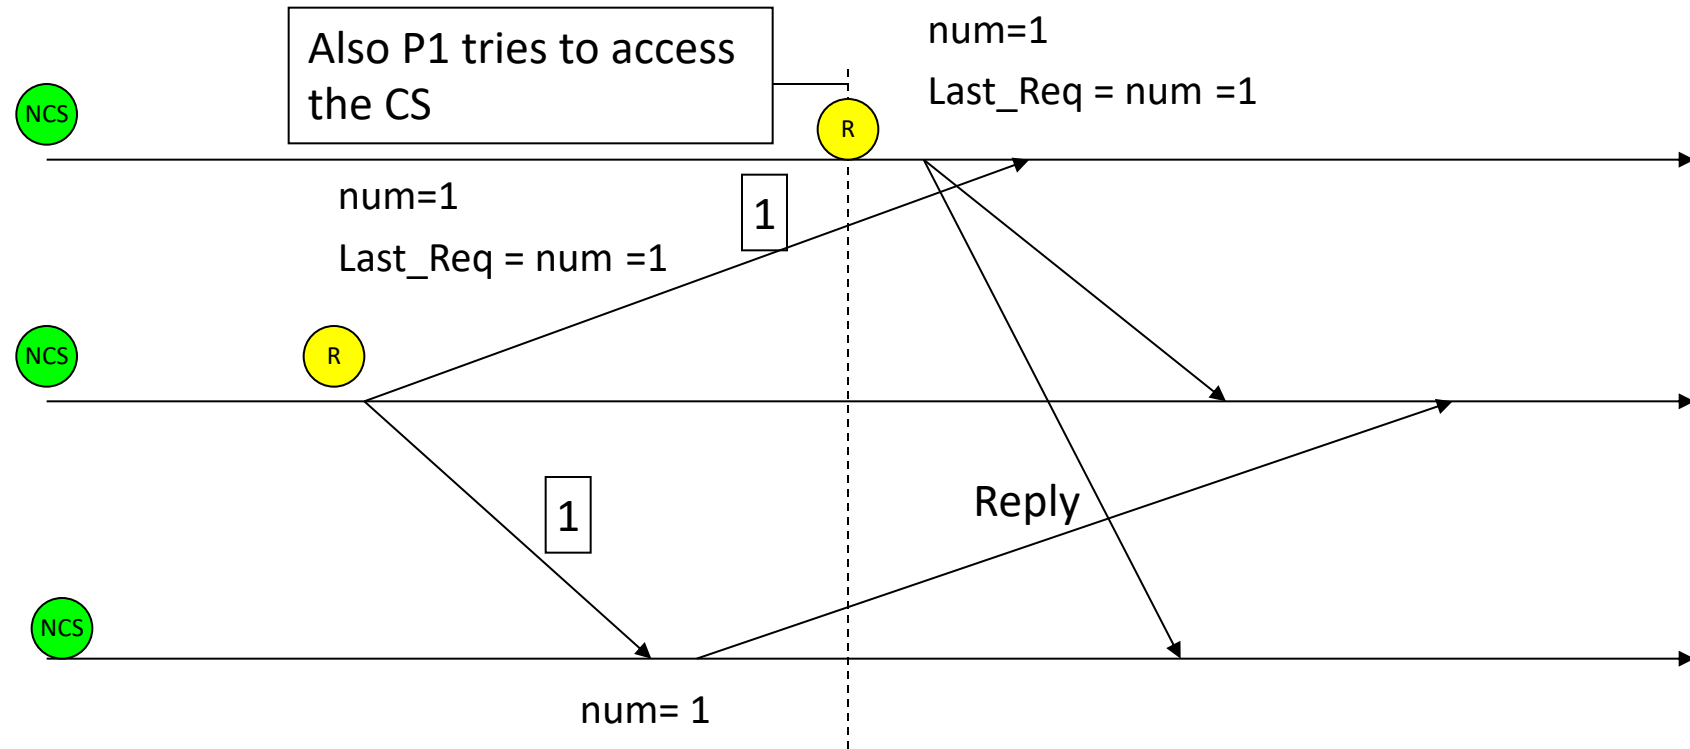
Ricart-Agrawala's algorithm: example



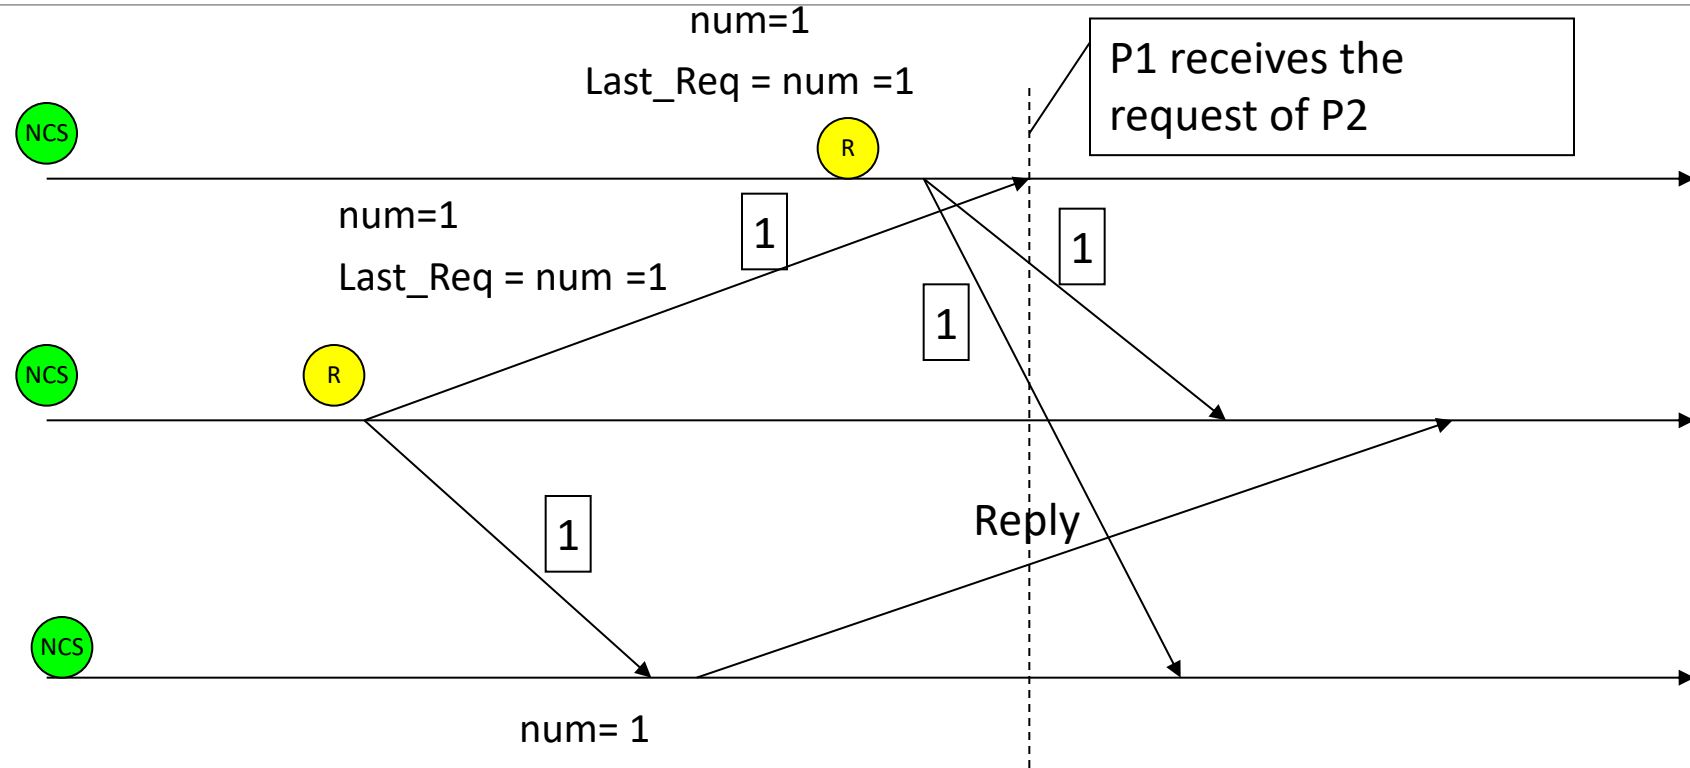
Ricart-Agrawala's algorithm: example



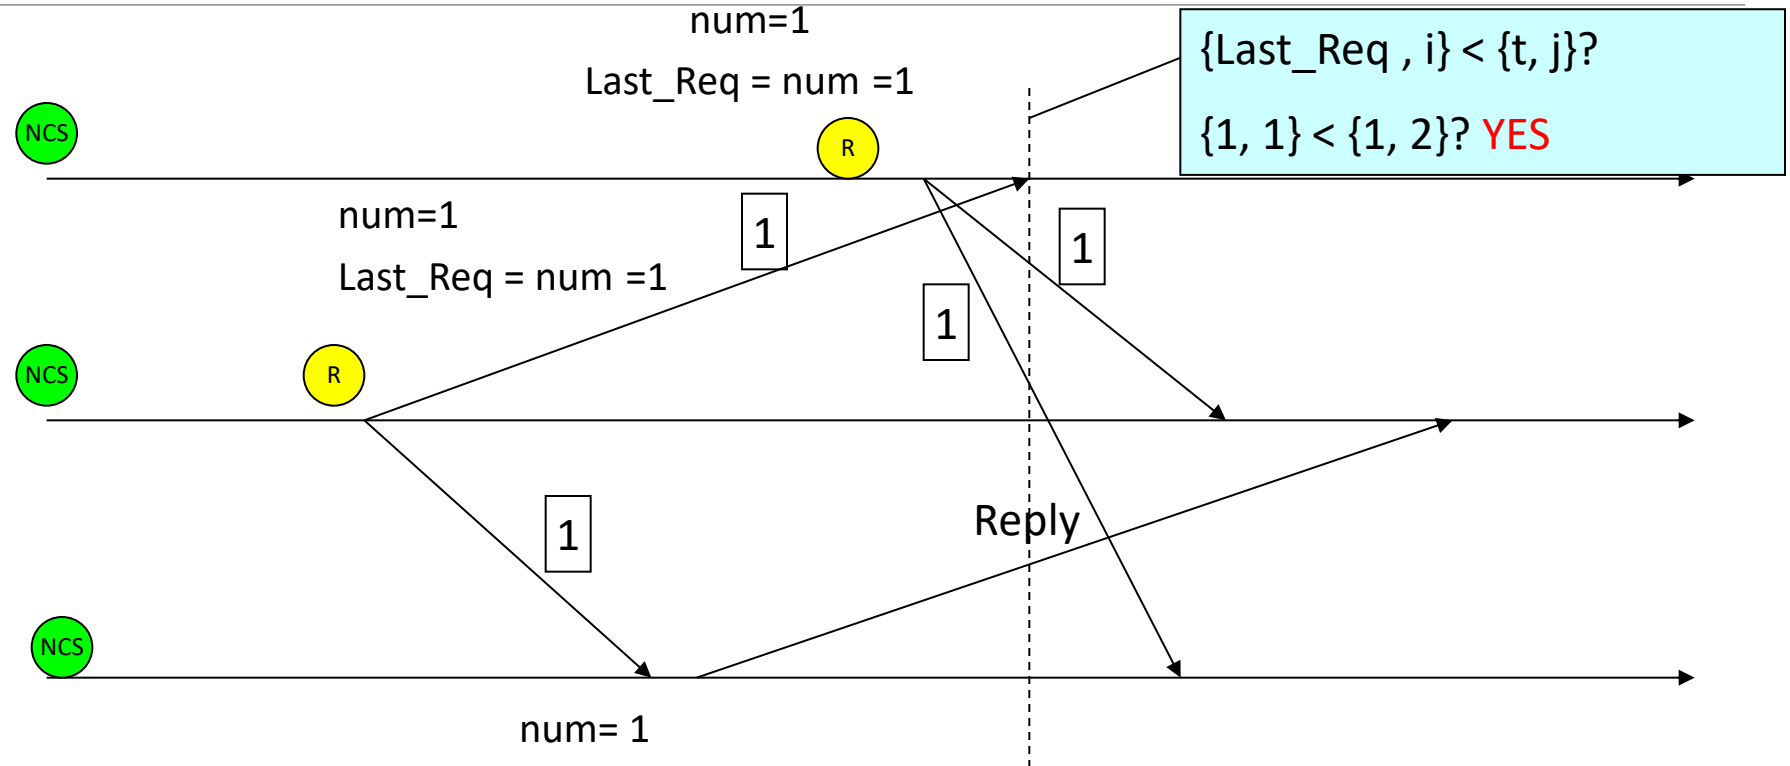
Ricart-Agrawala's algorithm: example



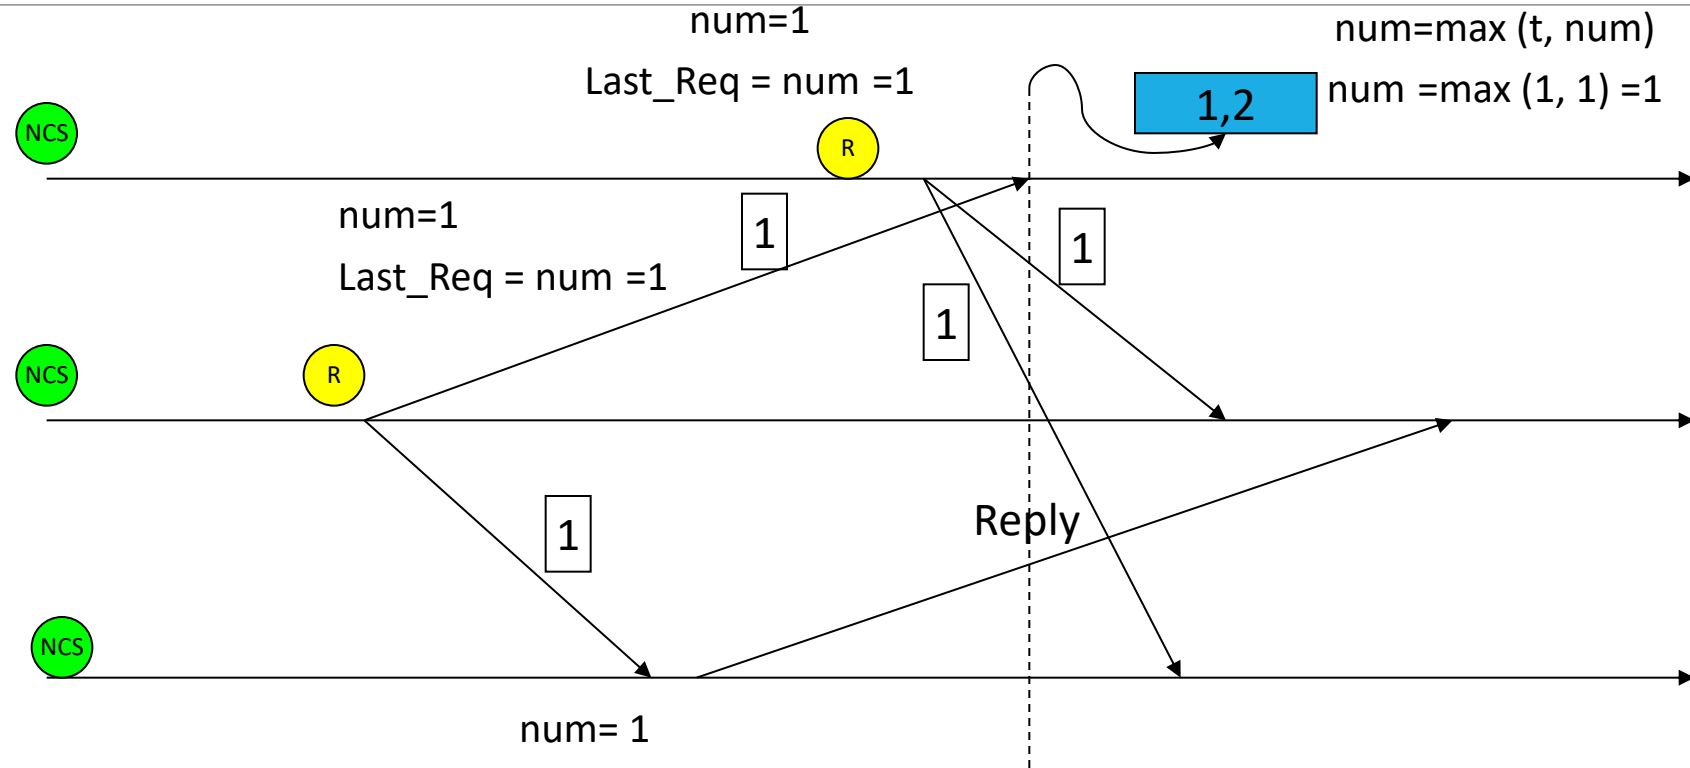
Ricart-Agrawala's algorithm: example



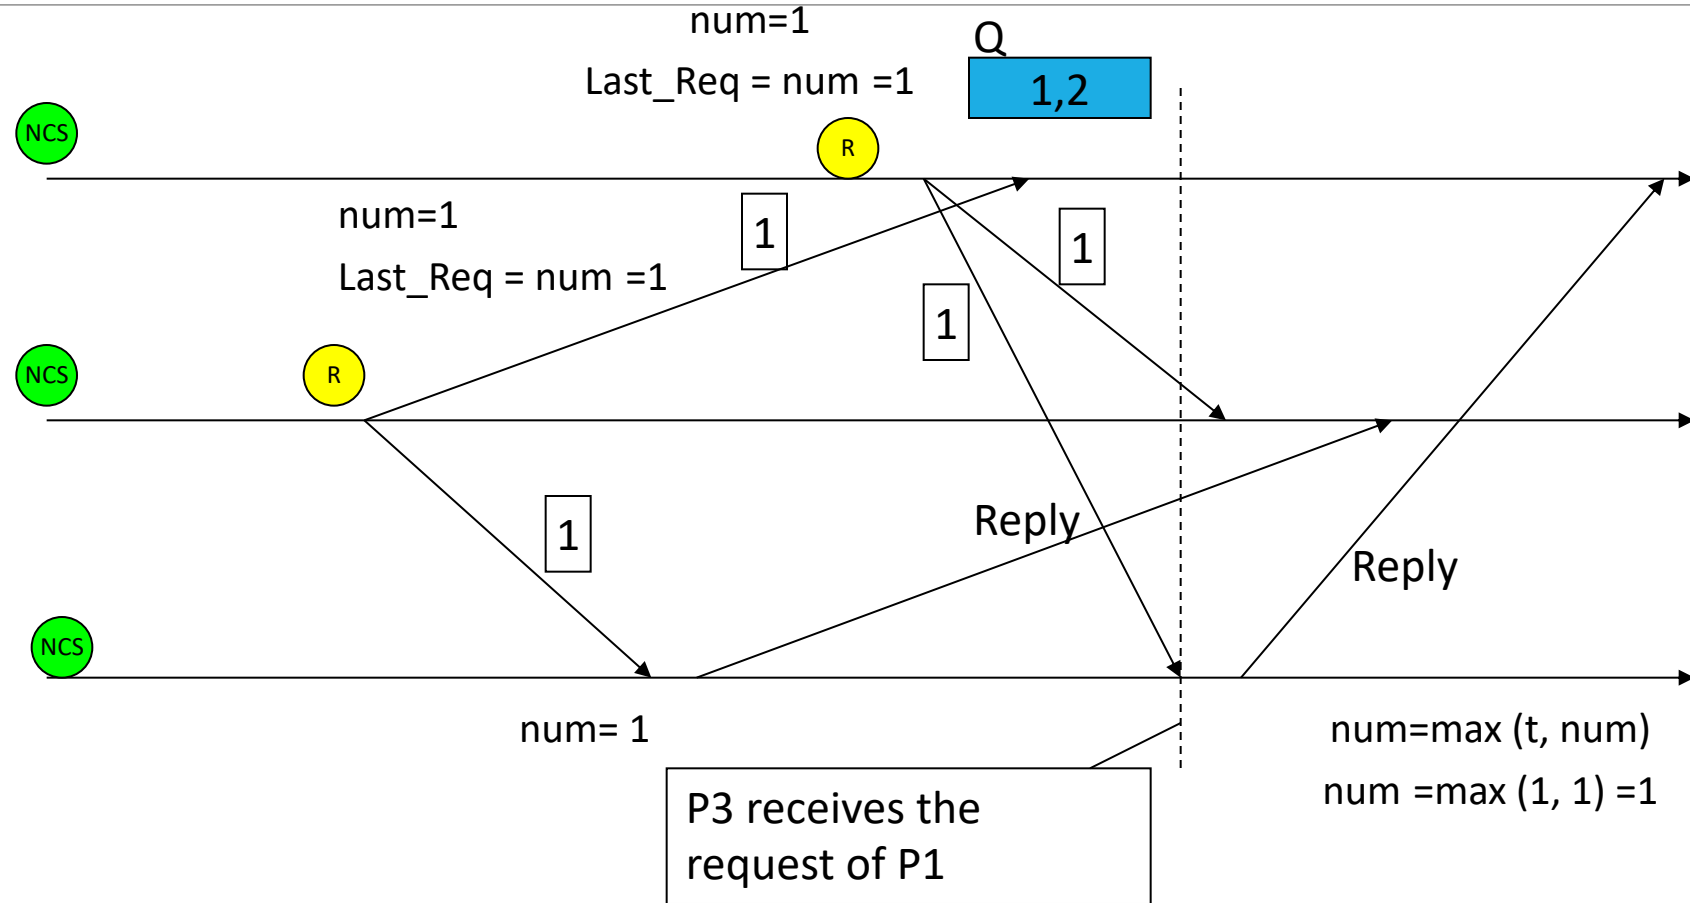
Ricart-Agrawala's algorithm: example



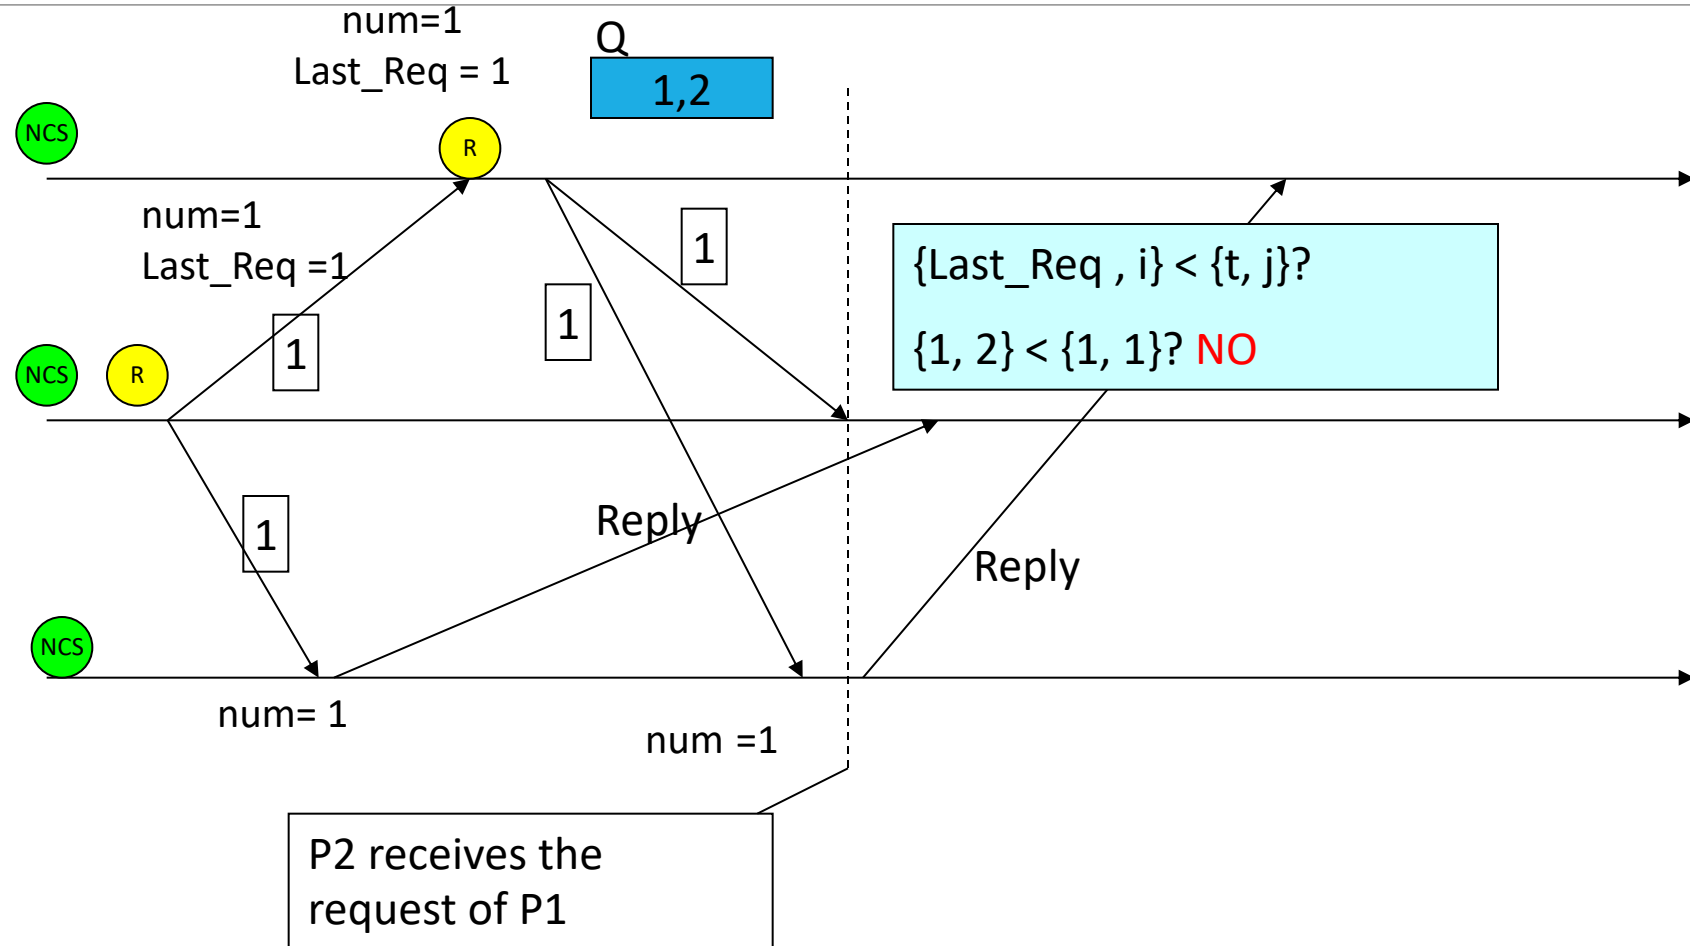
Ricart-Agrawala's algorithm: example




Ricart-Agrawala's algorithm: example



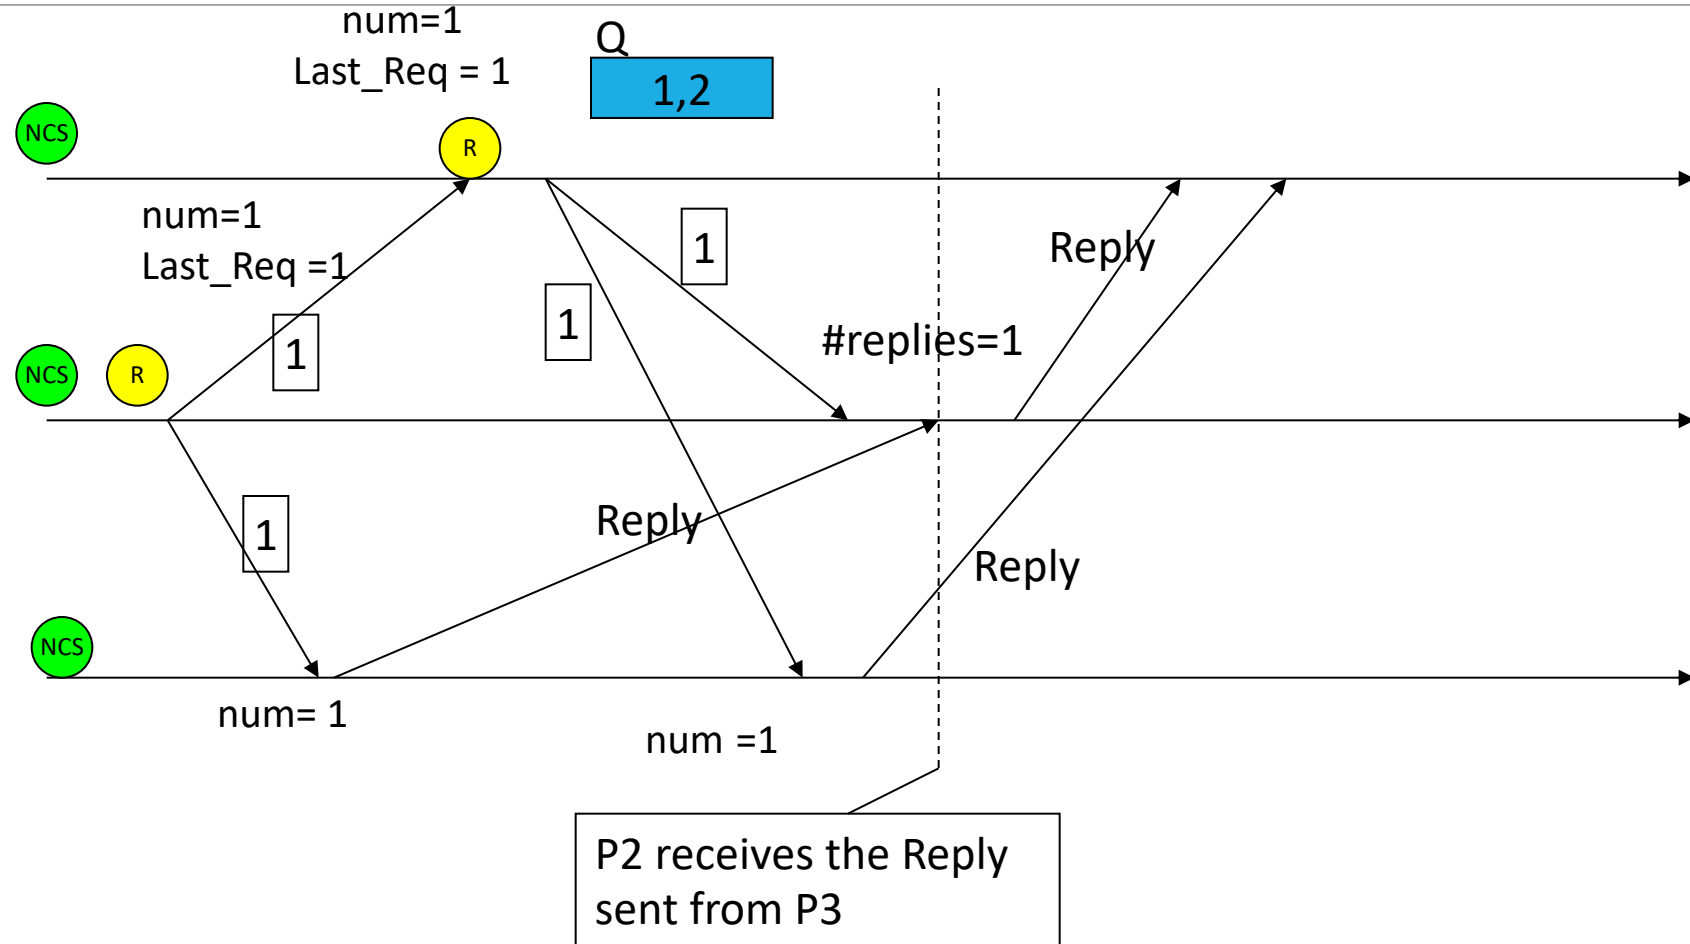
Ricart-Agrawala's algorithm: example



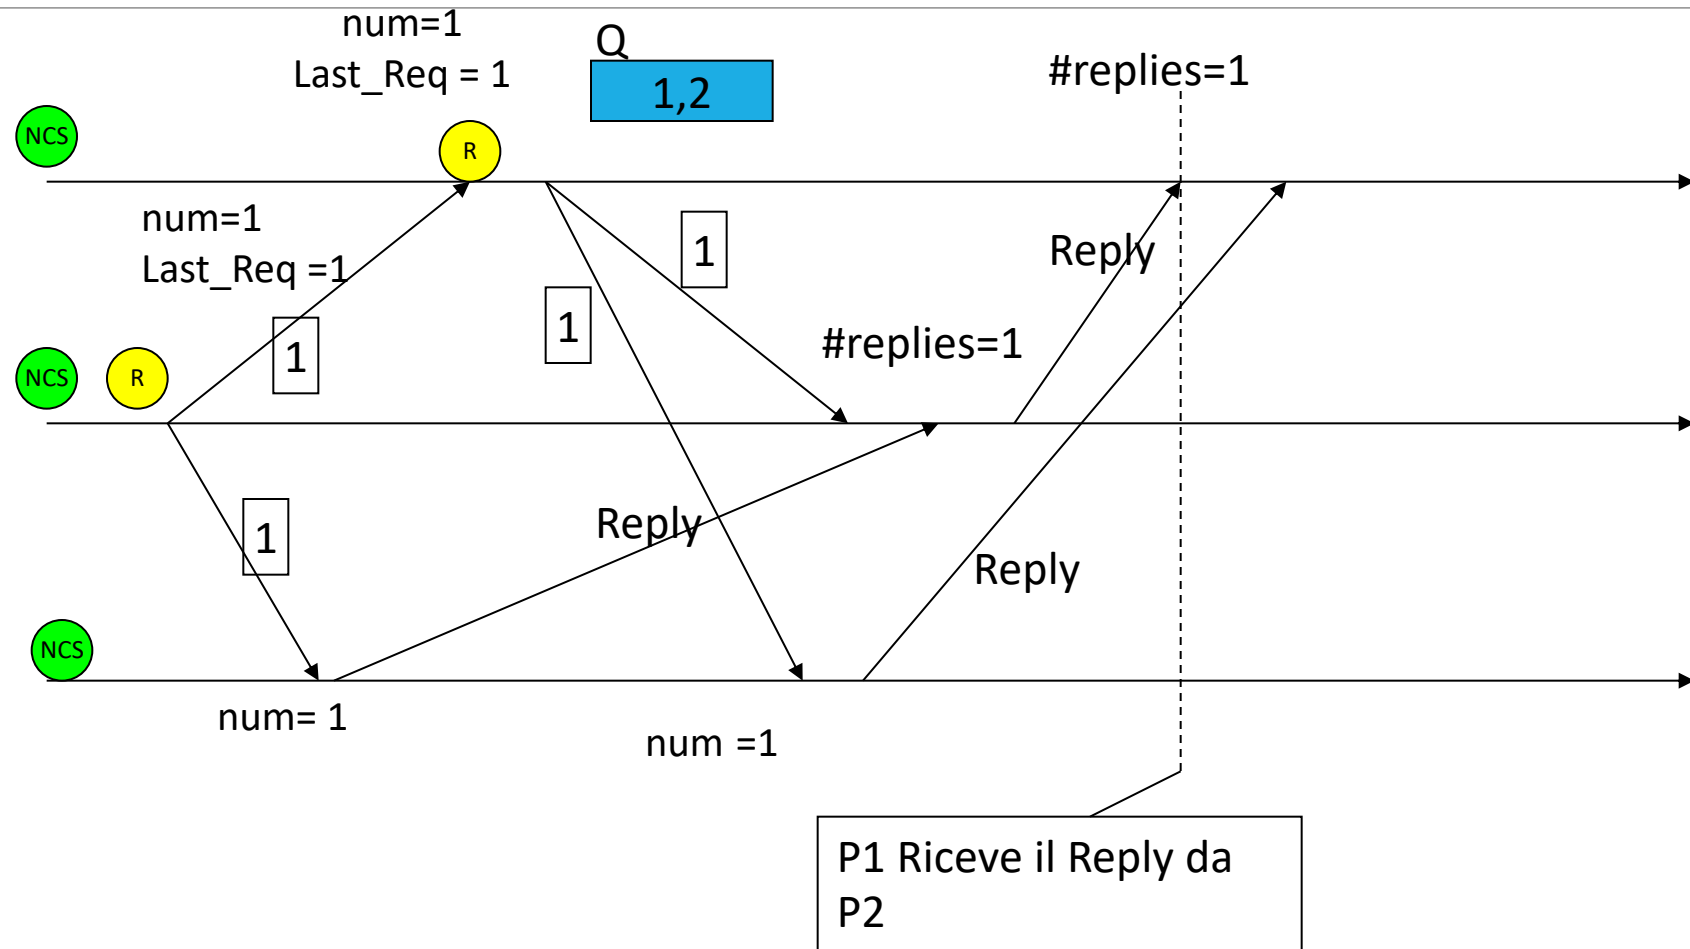
Ricart-Agrawala's algorithm: example




Ricart-Agrawala's algorithm: example



Ricart-Agrawala's algorithm: example



Ricart-Agrawala's algorithm: example



Ricart-Agrawala's algorithm: example

