

# Distributed Systems

## Master of Science in Engineering in Computer Science

AA 2019/2020

---

LECTURE 4: FAILURE DETECTION

# Recap on Timing Assumptions

---

## Synchronous

- timing assumptions are explicit either on
  - Bounds on process executions and communication channels, or
  - Existence of a common global clock, or
  - Both

## Asynchronous

- there are no timing assumptions

# Recap on Timing Assumptions

---

Partial synchrony requires abstract timing assumptions (after an unknown time  $t$  the system becomes synchronous)

Two choices:

1. Put assumption on the system model (including links and processes)
2. Create a separate abstractions that encapsulates those timing assumptions

Note: manipulating time inside a protocol/algorithm is complex and the correctness proof may become very involved and sometimes prone to errors

# Failure Detector Abstraction

---

Software module to be used together with process and link abstractions

It encapsulates timing assumptions of a either partially synchronous or fully synchronous system

The stronger are the timing assumption, the more accurate the information provided by a failure detector will be.

Described by two properties:

- Accuracy (informally is the ability to avoid mistakes in the detection)
- Completeness (informally is ability to detect all failures)

# Perfect Failure detectors (P)

---

## System model

- synchronous system
- crash failures

Using a own clock and the bounds of the synchrony model, a process can infer if another process has crashed

# Perfect failure detectors (P)

## Specification

---

---

**Module 2.6:** Interface and properties of the perfect failure detector

---

**Module:**

**Name:** PerfectFailureDetector, **instance**  $\mathcal{P}$ .

**Events:**


**Indication:**  $\langle \mathcal{P}, \text{Crash} \mid p \rangle$ : Detects that process  $p$  has crashed.

**Properties:**

**PFD1:** *Strong completeness:* Eventually, every process that crashes is permanently detected by every correct process.

**PFD2:** *Strong accuracy:* If a process  $p$  is detected by any process, then  $p$  has crashed.

---



# Perfect failure detectors (P) Implementation

---

**Algorithm 2.5:** Exclude on Timeout

---

**Implements:**

PerfectFailureDetector, **instance**  $\mathcal{P}$ .

**Uses:**

PerfectPointToPointLinks, **instance**  $pl$ .

**upon event**  $\langle \mathcal{P}, \text{Init} \rangle$  **do**

$alive := \Pi$ ;  
 $detected := \emptyset$ ;  
 $starttimer(\Delta)$ ;

**upon event**  $\langle \text{Timeout} \rangle$  **do**

**forall**  $p \in \Pi$  **do**  
    **if**  $(p \notin alive) \wedge (p \notin detected)$  **then**  
         $detected := detected \cup \{p\}$ ;  
        **trigger**  $\langle \mathcal{P}, \text{Crash} \mid p \rangle$ ;  
        **trigger**  $\langle pl, \text{Send} \mid p, [\text{HEARTBEATREQUEST}] \rangle$ ;  
     $alive := \emptyset$ ;  
     $starttimer(\Delta)$ ;

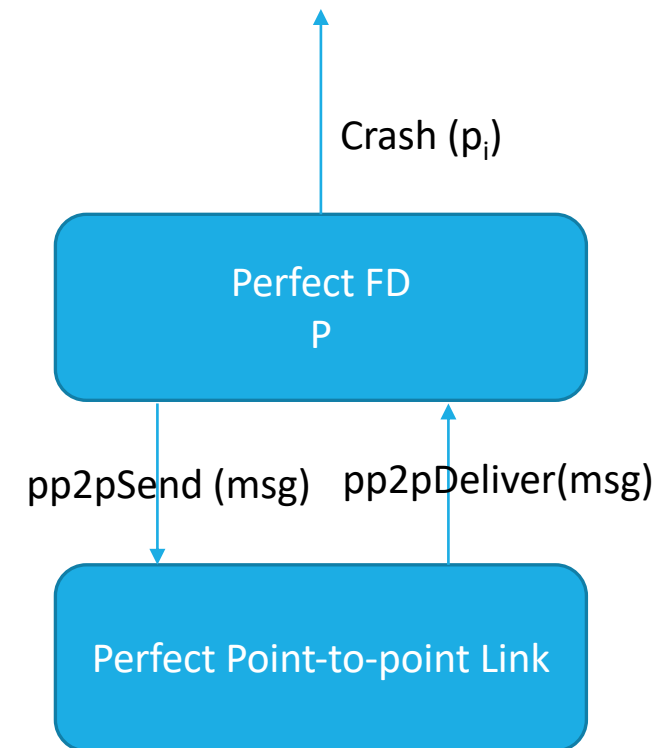
**upon event**  $\langle pl, \text{Deliver} \mid q, [\text{HEARTBEATREQUEST}] \rangle$  **do**

**trigger**  $\langle pl, \text{Send} \mid q, [\text{HEARTBEATREPLY}] \rangle$ ;

**upon event**  $\langle pl, \text{Deliver} \mid p, [\text{HEARTBEATREPLY}] \rangle$  **do**

$alive := alive \cup \{p\}$ ;

---



# Correctness

---

- To prove the correctness we have to prove that both Strong Completeness and Strong Accuracy are satisfied
- What if links are fair loss?
- What if we select a timeout too long?
- What if we select a timeout too short?



# Eventually perfect failure detectors ( $\diamond P$ )

---

## System model

- partial synchrony
- Crash failures
- Perfect point-to-point links

There is a (unknown) time  $t$  after that crashes can be accurately detected

Before  $t$  the systems behaves as an asynchronous one

The failure detector makes mistake in that periods assuming correct processes as crashed.

The notion of detection becomes suspicious

# Eventually perfect failure detectors ( $\diamond P$ ) Specification

---

---

**Module 2.8:** Interface and properties of the eventually perfect failure detector

---

**Module:**

**Name:** EventuallyPerfectFailureDetector, **instance**  $\diamond P$ .

**Events:**

**Indication:**  $\langle \diamond P, \text{Suspect} \mid p \rangle$ : Notifies that process  $p$  is suspected to have crashed.


**Indication:**  $\langle \diamond P, \text{Restore} \mid p \rangle$ : Notifies that process  $p$  is not suspected anymore.

**Properties:**

**EPFD1:** *Strong completeness:* Eventually, every process that crashes is permanently suspected by every correct process.

**EPFD2:** *Eventual strong accuracy:* Eventually, no correct process is suspected by any correct process.

---



# Basic constructions rules of an eventually perfect FD

---

- Use timeouts to suspect processes that did not sent expected messages
- A suspect may be wrong. A process  $p$  may suspect another one  $q$  because the chosen timeout was too short
- $\diamond$   $P$  is ready to reverse its judgment as soon as it receives a message from  $q$  (updating also the timeout value)
- If  $q$  has actually crashed,  $p$  does not change its judgment anymore.

# Eventually perfect failure detectors ( $\diamond P$ ) Implementation

**Algorithm 2.7:** Increasing Timeout

**Implements:**

EventuallyPerfectFailureDetector, **instance**  $\diamond P$ .

**Uses:**

PerfectPointToPointLinks, **instance**  $pl$ .

**upon event**  $\langle \diamond P, \text{Init} \rangle$  **do**

$alive := \Pi$ ;  
 $suspected := \emptyset$ ;  
 $delay := \Delta$ ;  
 $starttimer(delay)$ ;

**upon event**  $\langle \text{Timeout} \rangle$  **do**

**if**  $alive \cap suspected \neq \emptyset$  **then**  
 $delay := delay + \Delta$ ;

**forall**  $p \in \Pi$  **do**

**if**  $(p \notin alive) \wedge (p \notin suspected)$  **then**

$suspected := suspected \cup \{p\}$ ;

**trigger**  $\langle \diamond P, \text{Suspect} \mid p \rangle$ ;

**else if**  $(p \in alive) \wedge (p \in suspected)$  **then**

$suspected := suspected \setminus \{p\}$ ;

**trigger**  $\langle \diamond P, \text{Restore} \mid p \rangle$ ;

**trigger**  $\langle pl, \text{Send} \mid p, [\text{HEARTBEATREQUEST}] \rangle$ ;

$alive := \emptyset$ ;

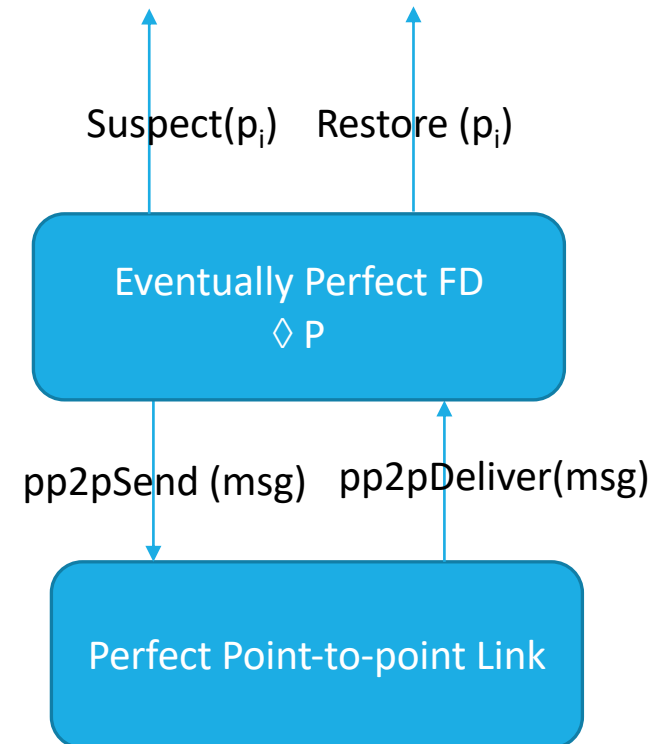
$starttimer(delay)$ ;

**upon event**  $\langle pl, \text{Deliver} \mid q, [\text{HEARTBEATREQUEST}] \rangle$  **do**

**trigger**  $\langle pl, \text{Send} \mid q, [\text{HEARTBEATREPLY}] \rangle$ ;

**upon event**  $\langle pl, \text{Deliver} \mid p, [\text{HEARTBEATREPLY}] \rangle$  **do**

$alive := alive \cup \{p\}$ ;



# Correctness

---

**Strong completeness.** If a process crashes, it will stop to send messages. Therefore the process will be suspected by any correct process and no process will revise the judgement.

**Eventual strong accuracy.** After time  $T$  the system becomes synchronous. i.e., after that time a message sent by a correct process  $p$  to another one  $q$  will be delivered within a bounded time. If  $p$  was wrongly suspected by  $q$ , then  $q$  will revise its suspicious.

# References

---

C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

- Chapter 2 – from Section 2.6.1 to Section 2.6.5